

# Weekly Summary Template

Brady Miller

## Table of contents

Thursday, April 6 . . . . .	1
-----------------------------	---

## Thursday, April 6

### ! TIL

Include a *very brief* summary of what you learnt in this class here.  
Today, I learnt the following concepts in class:

1. Luz Hyperparameters
2. Luz setup and fit process
3. Luz validation and metrics that we can use

```
packages <- c(
  "dplyr",
  "readr",
  "tidyr",
  "purrr",
  "stringr",
  "corrplot",
  "car",
  "caret",
  "torch",
  "nnet",
  "broom",
  "torch",
  "torchvision",
  "e1071",
```

```
"glmnet",  
"nnet",  
"rpart",  
"ISLR2",  
'luz',  
'torchvision'  
)  
  
renv::install(packages)
```

```
Installing dplyr [1.1.1] ...  
  OK [linked cache in 2.4 milliseconds]  
Installing readr [2.1.4] ...  
  OK [linked cache in 2.7 milliseconds]  
Installing purrr [1.0.1] ...  
  OK [linked cache in 2.6 milliseconds]  
Installing stringr [1.5.0] ...  
  OK [linked cache in 3.1 milliseconds]  
Installing tidyr [1.3.0] ...  
  OK [linked cache in 3.4 milliseconds]  
Installing corrplot [0.92] ...  
  OK [linked cache in 2.4 milliseconds]  
Installing nnet [7.3-18] ...  
  OK [linked cache in 3.1 milliseconds]  
Installing broom [1.0.4] ...  
  OK [linked cache in 3.1 milliseconds]  
Installing car [3.1-2] ...  
  OK [linked cache in 2.8 milliseconds]  
Installing e1071 [1.7-13] ...  
  OK [linked cache in 2.4 milliseconds]  
Installing rpart [4.1.19] ...  
  OK [linked cache in 3.2 milliseconds]  
Installing caret [6.0-94] ...  
  OK [linked cache in 5 milliseconds]  
Installing torch [0.9.1] ...  
  OK [linked cache in 2.3 milliseconds]  
Installing torchvision [0.5.0] ...  
  OK [linked cache in 2.4 milliseconds]  
Installing glmnet [4.1-7] ...  
  OK [linked cache in 2.5 milliseconds]  
Installing ISLR2 [1.3-2] ...  
  OK [linked cache in 2.5 milliseconds]
```

```
Installing luz [0.3.1] ...  
OK [linked cache in 2.4 milliseconds]
```

```
sapply(packages, require, character.only=T)
```

```
Loading required package: dplyr
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
Loading required package: readr
```

```
Loading required package: tidyr
```

```
Loading required package: purrr
```

```
Loading required package: stringr
```

```
Loading required package: corrplot
```

```
corrplot 0.92 loaded
```

```
Loading required package: car
```

```
Loading required package: carData
```

```
Attaching package: 'car'
```

```
The following object is masked from 'package:purrr':
```

```
some
```

The following object is masked from 'package:dplyr':

recode

Loading required package: caret

Loading required package: ggplot2

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

Loading required package: torch

Loading required package: nnet

Loading required package: broom

Loading required package: torchvision

Loading required package: e1071

Loading required package: glmnet

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-7

Loading required package: rpart

Loading required package: ISLR2


Loading required package: luz

dplyr	readr	tidyr	purrr	stringr	corrplot
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
car	caret	torch	nnet	broom	torch
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
torchvision	e1071	glmnet	nnet	rpart	ISLR2
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
luz	torchvision				
TRUE	TRUE				

Creating a dataset that we can use with luz and neural network to make predictions with

```
ex <- \(x) ifelse(
  ((abs(x[1]) + 0.05 * rnorm(1) > 0.50 & abs(x[2]) + 0.05 * rnorm(1) > 0.50)) |
  ((abs(x[1]) + 0.05 * rnorm(1) < 0.25 & abs(x[2]) + 0.05 * rnorm(1) < 0.25)), 1, 0
)

n <- 300
X <- t(replicate(n, 2 * runif(2) - 1))
y <- apply(X, 1, ex) %>% as.factor()
col <- ifelse(y == 0, 'blue', 'red')
df <- data.frame(y = y, x1 = X[,1], x2 = X[,2])
model <- glm(y ~ x1 + x2, df, family = binomial())
plot(df$x1, df$x2, col = col, pch = 19)
```



```
xnew <- cbind(
  rep(seq(-1.1, 1.1, length.out = 50), 50),
  rep(seq(-1.1, 1.1, length.out = 50), each = 50)
)

df_new = data.frame(x1 = xnew[,1], x2 = xnew[,2])
```

## Luz Hyperparameters

We can use hyperparameters to specify the inputs for the neural network so we can create our own inputs/dimensions for the hidden layers

```
nn_model <- nn_module(
  initialize = function(p, q1, q2, q3) {
    self$f <- nn_linear(p, q1)
    self$g <- nn_linear(q1, q2)
    self$h <- nn_linear(q2, q3)
    self$i <- nn_linear(q3, 1)
    self$a <- nn_relu()
    self$s <- nn_sigmoid()
  },
```

```

forward = function(x) {
  x %>%
    self$f() %>%
    self$a() %>%
    self$g() %>%
    self$a() %>%
    self$h() %>%
    self$a() %>%
    self$i() %>%
    self$s()
}
)

```

## Luz Setup & Fit

In this next section, I will show how to use luz to setup and fit the nn model

```

nn_model %>%
  setup(
    loss = nn_bce_loss(),
    optimizer = optim_adam
  )

```

<luz\_module\_generator>

\*Setup function takes in 2 mandatory arguments -> the loss function and the optimizer (used for minimizing loss) 1. optimizer could be optim\_adam, optim\_rmsprop, ...

The code above is equivalent to...

```

F <- hh2_module()
optimizer <- optim_adam(F$parameters, lr = 0.05)
epochs <- 1000

for (i in 1:epochs){
  loss <- nn_mse_loss()(F(X_tensor), y_tensor)
  optimizer$zero_grad()
  loss$backward()
  optimizer$step()
}

```

Other things we may want to specify

1. Epochs
2. Gradient descent step
3. x,y as tensors
4. Learning rate
5. what p & q are

These things listed above are now added/specified in the code below

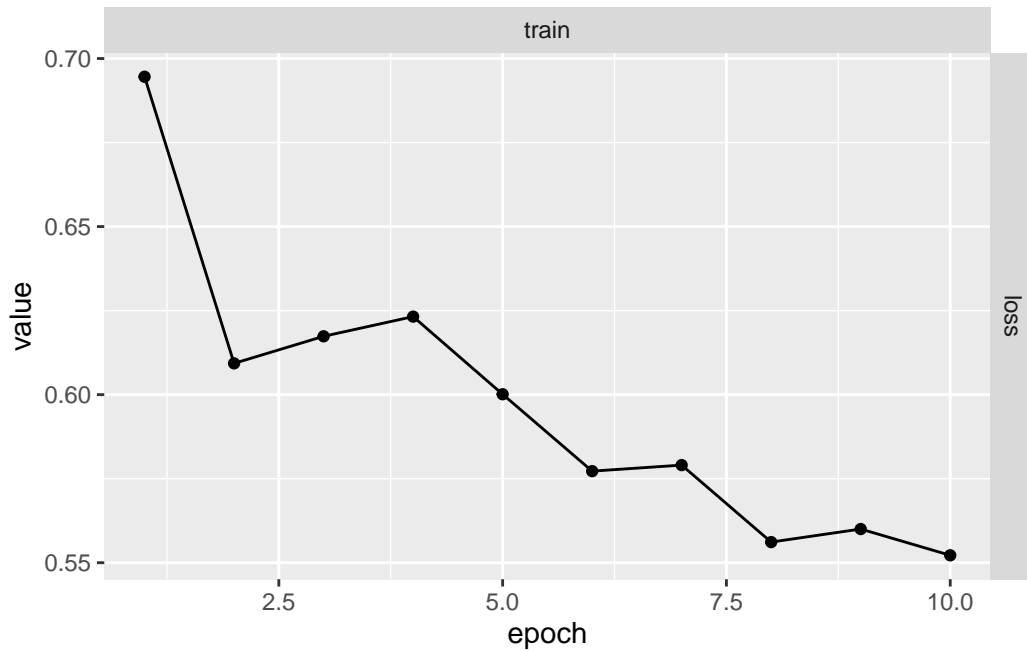
```
fit_nn <- nn_model %>%
  setup(
    loss = nn_bce_loss(),
    optimizer = optim_adam
  ) %>%
  set_hparams(p = 2, q1 = 5, q2 = 7, q3 = 5) %>%
  set_opt_hparams(lr = 0.02) %>%
  # Fit the neural network
  # Have to change formatting b/c torch can only read matrices, not data frames
  fit(
    data = list(
      as.matrix(df[, -1]),
      as.numeric(df[, 1]) - 1
    ),
    epochs = 10,
    verbose = TRUE
  )
```

```
Epoch 1/10
Train metrics: Loss: 0.6946
Epoch 2/10
Train metrics: Loss: 0.6093
Epoch 3/10
Train metrics: Loss: 0.6174
Epoch 4/10
Train metrics: Loss: 0.6232
Epoch 5/10
Train metrics: Loss: 0.6001
Epoch 6/10
Train metrics: Loss: 0.5772
Epoch 7/10
Train metrics: Loss: 0.579
Epoch 8/10
Train metrics: Loss: 0.5561
```



```
Epoch 9/10
Train metrics: Loss: 0.56
Epoch 10/10
Train metrics: Loss: 0.5522
```

```
# plots change in loss for the epochs specified
plot(fit_nn)
```



Based on the plot, we can see that loss does decrease a little bit, but it jumps around, both increasing and decreasing after the initial decrease.

The output of the Luz allows you to use the predict function

```
predict(fit_nn, xnew)
```

```
torch_tensor
0.7034
0.6928
0.6820
0.6710
0.6599
```

```

0.6485
0.6370
0.6253
0.6135
0.5735
0.5256
0.4772
0.4293
0.3827
0.3381
0.2962
0.2575
0.2264
0.2026
0.1806
0.1698
0.1639
0.1653
0.1666
0.1679
0.1693
0.1706
0.1720
0.1734
0.1747
... [the output was truncated (use n=-1 to disable)]
[ CPUFloatType{2500,1} ]

```

```

predict(fit_nn, cbind(rnorm(10), rnorm(10))) %>% as.array

```

```

      [,1]
[1,] 0.5580807
[2,] 0.4615246
[3,] 0.5967558
[4,] 0.1521935
[5,] 0.5666237
[6,] 0.6143605
[7,] 0.6944618
[8,] 0.4637559
[9,] 0.1690233
[10,] 0.4269857

```

## Luz Validation Data and Metrics

Randomly selecting the indices of 23 rows in the data frame without replacement

```
test_ind <- sample(1:nrow(df), 23, replace = FALSE)
```

Using the luz code we created to validate the data

```
fit_nn <- nn_model %>%  
  setup(  
    loss = nn_bce_loss(),  
    optimizer = optim_adam  
  ) %>%  
  set_hparams(p = 2, q1 = 5, q2 = 7, q3 = 5) %>%  
  set_opt_hparams(lr = 0.02) %>%  
  fit(  
    data = list(  
      as.matrix(df[-test_ind,-1]),  
      as.numeric(df[-test_ind,1]) - 1  
    ),  
    valid_data = list(  
      as.matrix(df[+test_ind, -1]),  
      as.numeric(df[+test_ind,1]) - 1  
    ),  
    epochs = 10,  
    verbose = TRUE  
  )
```

Epoch 1/10

Train metrics: Loss: 0.6286

Valid metrics: Loss: 0.5894

Epoch 2/10

Train metrics: Loss: 0.6188

Valid metrics: Loss: 0.5829

Epoch 3/10

Train metrics: Loss: 0.6164

Valid metrics: Loss: 0.5807

Epoch 4/10

Train metrics: Loss: 0.6137


Valid metrics: Loss: 0.5788

Epoch 5/10

Train metrics: Loss: 0.6017

```
Valid metrics: Loss: 0.5745
Epoch 6/10
Train metrics: Loss: 0.6021
Valid metrics: Loss: 0.5709
Epoch 7/10
Train metrics: Loss: 0.5829
Valid metrics: Loss: 0.5616
Epoch 8/10
Train metrics: Loss: 0.5568
Valid metrics: Loss: 0.5516
Epoch 9/10
Train metrics: Loss: 0.5258
Valid metrics: Loss: 0.5232
Epoch 10/10
Train metrics: Loss: 0.5072
Valid metrics: Loss: 0.5057
```

```
plot(fit_nm)
```



From this plot we can see that the loss generally decreases of for the training data and has a greater loss value than the test data throughout the epochs, but the test data switches between increasing and decreasing as the number of epochs increases.

- Luz has built in metrics(ex. accuracy, mse, ...) some of which are shown below

```
predicted <- torch_randn(100)
expected <- torch_randn(100)
metric <- luz_metric_binary_accuracy()

metric <- metric$new()
metric$update(predicted, expected)
metric$compute()
```

```
[1] 0
```

This time we are specifying the metrics we want to include, in the fitting of the neural network model

```
fit_nn <- nn_model %>%
  setup(
    loss = nn_bce_loss(),
    optimizer = optim_adam,
    # specifying metrics we want to use
    metrics = list(
      luz_metric_binary_accuracy(),
      luz_metric_binary_auroc()
    )
  ) %>%
  set_hparams(p = 2, q1 = 5, q2 = 7, q3 = 5) %>%
  set_opt_hparams(lr = 0.02) %>%
  fit(
    data = list(
      as.matrix(df[-test_ind,-1]),
      as.numeric(df[-test_ind,1]) - 1
    ),
    valid_data = list(
      as.matrix(df[+test_ind, -1]),
      as.numeric(df[+test_ind,1]) - 1
    ),
    epochs = 50,
    verbose = TRUE
  )
```

Epoch 1/50

Train metrics: Loss: 0.7107 - Acc: 13.3285 - AUC: 0.4212  
Valid metrics: Loss: 0.6515 - Acc: 17 - AUC: 0  
Epoch 2/50  
Train metrics: Loss: 0.6348 - Acc: 21.3538 - AUC: 0.5155  
Valid metrics: Loss: 0.5881 - Acc: 17 - AUC: 0.1667  
Epoch 3/50  
Train metrics: Loss: 0.606 - Acc: 21.2744 - AUC: 0.5583  
Valid metrics: Loss: 0.5766 - Acc: 17 - AUC: 0.1667  
Epoch 4/50  
Train metrics: Loss: 0.6029 - Acc: 21.4729 - AUC: 0.5408  
Valid metrics: Loss: 0.5628 - Acc: 17 - AUC: 0.1667  
Epoch 5/50  
Train metrics: Loss: 0.5841 - Acc: 20.9964 - AUC: 0.4981  
Valid metrics: Loss: 0.5592 - Acc: 16.5217 - AUC: 0.1667  
Epoch 6/50  
Train metrics: Loss: 0.5729 - Acc: 20.5415 - AUC: 0.5159  
Valid metrics: Loss: 0.5515 - Acc: 17 - AUC: 0.1667  
Epoch 7/50  
Train metrics: Loss: 0.5578 - Acc: 20.5884 - AUC: 0.5284  
Valid metrics: Loss: 0.5412 - Acc: 16.5217 - AUC: 0.1667  
Epoch 8/50  
Train metrics: Loss: 0.5397 - Acc: 20.4585 - AUC: 0.5642  
Valid metrics: Loss: 0.5426 - Acc: 16.5217 - AUC: 0.1667  
Epoch 9/50  
Train metrics: Loss: 0.5394 - Acc: 20.1552 - AUC: 0.5805  
Valid metrics: Loss: 0.5356 - Acc: 16.5217 - AUC: 0.1667  
Epoch 10/50  
Train metrics: Loss: 0.5414 - Acc: 20.0866 - AUC: 0.5892  
Valid metrics: Loss: 0.5337 - Acc: 16.5217 - AUC: 0.1667  
Epoch 11/50  
Train metrics: Loss: 0.5285 - Acc: 20.3827 - AUC: 0.5969  
Valid metrics: Loss: 0.5315 - Acc: 16.5217 - AUC: 0.1667  
Epoch 12/50  
Train metrics: Loss: 0.5386 - Acc: 19.8267 - AUC: 0.5773  
Valid metrics: Loss: 0.53 - Acc: 16.5217 - AUC: 0.1667  
Epoch 13/50  
Train metrics: Loss: 0.5225 - Acc: 20.0722 - AUC: 0.5574  
Valid metrics: Loss: 0.526 - Acc: 16.5217 - AUC: 0.1667  
Epoch 14/50  
Train metrics: Loss: 0.5239 - Acc: 19.8989 - AUC: 0.5583  
Valid metrics: Loss: 0.5313 - Acc: 16.5217 - AUC: 0.1667  
Epoch 15/50  
Train metrics: Loss: 0.5251 - Acc: 20.1805 - AUC: 0.5858

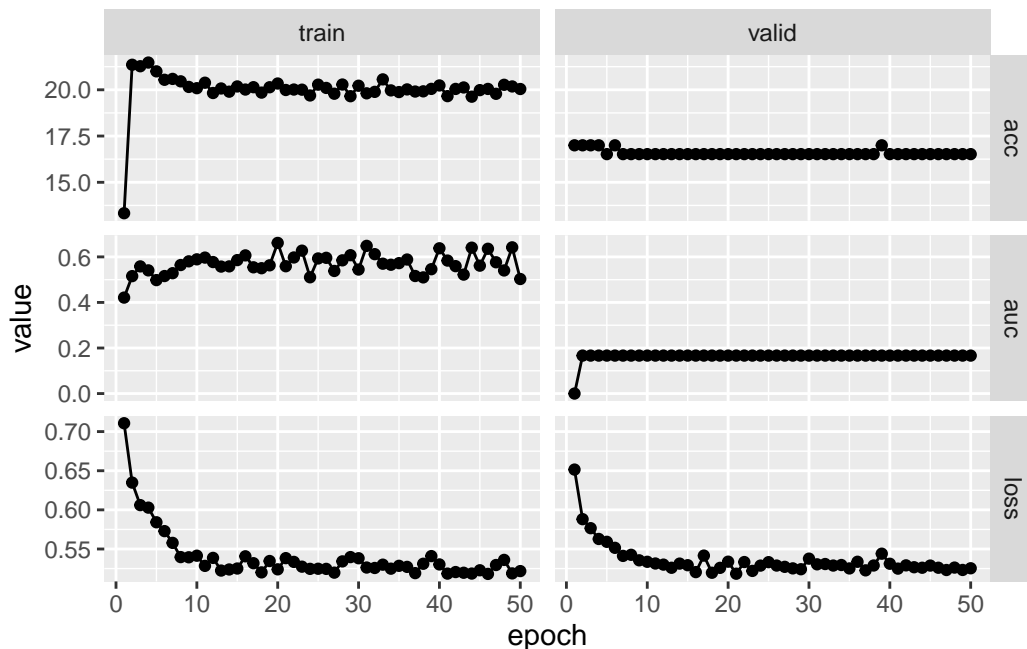
Valid metrics: Loss: 0.5293 - Acc: 16.5217 - AUC: 0.1667  
Epoch 16/50  
Train metrics: Loss: 0.5406 - Acc: 20.0144 - AUC: 0.6066  
Valid metrics: Loss: 0.5205 - Acc: 16.5217 - AUC: 0.1667  
Epoch 17/50  
Train metrics: Loss: 0.5319 - Acc: 20.1408 - AUC: 0.5543  
Valid metrics: Loss: 0.5416 - Acc: 16.5217 - AUC: 0.1667  
Epoch 18/50  
Train metrics: Loss: 0.5201 - Acc: 19.8484 - AUC: 0.5503  
Valid metrics: Loss: 0.5197 - Acc: 16.5217 - AUC: 0.1667  
Epoch 19/50  
Train metrics: Loss: 0.5346 - Acc: 20.1372 - AUC: 0.563  
Valid metrics: Loss: 0.5259 - Acc: 16.5217 - AUC: 0.1667  
Epoch 20/50  
Train metrics: Loss: 0.524 - Acc: 20.3394 - AUC: 0.6615  
Valid metrics: Loss: 0.5338 - Acc: 16.5217 - AUC: 0.1667  
Epoch 21/50  
Train metrics: Loss: 0.5383 - Acc: 19.9892 - AUC: 0.5589  
Valid metrics: Loss: 0.5184 - Acc: 16.5217 - AUC: 0.1667  
Epoch 22/50  
Train metrics: Loss: 0.5335 - Acc: 20.0144 - AUC: 0.5975  
Valid metrics: Loss: 0.5333 - Acc: 16.5217 - AUC: 0.1667  
Epoch 23/50  
Train metrics: Loss: 0.5274 - Acc: 20.0072 - AUC: 0.6273  
Valid metrics: Loss: 0.522 - Acc: 16.5217 - AUC: 0.1667  
Epoch 24/50  
Train metrics: Loss: 0.5246 - Acc: 19.6931 - AUC: 0.5105  
Valid metrics: Loss: 0.5288 - Acc: 16.5217 - AUC: 0.1667  
Epoch 25/50  
Train metrics: Loss: 0.5248 - Acc: 20.278 - AUC: 0.5932  
Valid metrics: Loss: 0.5332 - Acc: 16.5217 - AUC: 0.1667  
Epoch 26/50  
Train metrics: Loss: 0.5246 - Acc: 20.1011 - AUC: 0.596  
Valid metrics: Loss: 0.5289 - Acc: 16.5217 - AUC: 0.1667  
Epoch 27/50  
Train metrics: Loss: 0.5198 - Acc: 19.787 - AUC: 0.5382  
Valid metrics: Loss: 0.5272 - Acc: 16.5217 - AUC: 0.1667  
Epoch 28/50  
Train metrics: Loss: 0.5342 - Acc: 20.2852 - AUC: 0.5848  
Valid metrics: Loss: 0.5253 - Acc: 16.5217 - AUC: 0.1667  
Epoch 29/50  
Train metrics: Loss: 0.5397 - Acc: 19.6534 - AUC: 0.6074  
Valid metrics: Loss: 0.5242 - Acc: 16.5217 - AUC: 0.1667

Epoch 30/50  
Train metrics: Loss: 0.5382 - Acc: 20.2202 - AUC: 0.5444  
Valid metrics: Loss: 0.5377 - Acc: 16.5217 - AUC: 0.1667  
Epoch 31/50  
Train metrics: Loss: 0.5263 - Acc: 19.8087 - AUC: 0.6484  
Valid metrics: Loss: 0.5304 - Acc: 16.5217 - AUC: 0.1667  
Epoch 32/50  
Train metrics: Loss: 0.5258 - Acc: 19.8881 - AUC: 0.612  
Valid metrics: Loss: 0.5308 - Acc: 16.5217 - AUC: 0.1667  
Epoch 33/50  
Train metrics: Loss: 0.53 - Acc: 20.5632 - AUC: 0.5701  
Valid metrics: Loss: 0.5289 - Acc: 16.5217 - AUC: 0.1667  
Epoch 34/50  
Train metrics: Loss: 0.5249 - Acc: 19.9675 - AUC: 0.5655  
Valid metrics: Loss: 0.5295 - Acc: 16.5217 - AUC: 0.1667  
Epoch 35/50  
Train metrics: Loss: 0.5285 - Acc: 19.8773 - AUC: 0.5726  
Valid metrics: Loss: 0.5252 - Acc: 16.5217 - AUC: 0.1667  
Epoch 36/50  
Train metrics: Loss: 0.5272 - Acc: 20.0217 - AUC: 0.5885  
Valid metrics: Loss: 0.5337 - Acc: 16.5217 - AUC: 0.1667  
Epoch 37/50  
Train metrics: Loss: 0.5192 - Acc: 19.9061 - AUC: 0.5162  
Valid metrics: Loss: 0.5227 - Acc: 16.5217 - AUC: 0.1667  
Epoch 38/50  
Train metrics: Loss: 0.5311 - Acc: 19.917 - AUC: 0.51  
Valid metrics: Loss: 0.5289 - Acc: 16.5217 - AUC: 0.1667  
Epoch 39/50  
Train metrics: Loss: 0.5409 - Acc: 20.0505 - AUC: 0.5459  
Valid metrics: Loss: 0.544 - Acc: 17 - AUC: 0.1667  
Epoch 40/50  
Train metrics: Loss: 0.5303 - Acc: 20.2347 - AUC: 0.6377  
Valid metrics: Loss: 0.5312 - Acc: 16.5217 - AUC: 0.1667  
Epoch 41/50  
Train metrics: Loss: 0.5186 - Acc: 19.657 - AUC: 0.5838  
Valid metrics: Loss: 0.5249 - Acc: 16.5217 - AUC: 0.1667  
Epoch 42/50  
Train metrics: Loss: 0.5205 - Acc: 20.0469 - AUC: 0.5593  
Valid metrics: Loss: 0.5291 - Acc: 16.5217 - AUC: 0.1667  
Epoch 43/50  
Train metrics: Loss: 0.5197 - Acc: 20.1264 - AUC: 0.5222  
Valid metrics: Loss: 0.5267 - Acc: 16.5217 - AUC: 0.1667  
Epoch 44/50



Train metrics: Loss: 0.5186 - Acc: 19.6282 - AUC: 0.6405  
 Valid metrics: Loss: 0.5261 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 45/50  
 Train metrics: Loss: 0.5228 - Acc: 19.9819 - AUC: 0.5615  
 Valid metrics: Loss: 0.5288 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 46/50  
 Train metrics: Loss: 0.5182 - Acc: 20.0469 - AUC: 0.6356  
 Valid metrics: Loss: 0.5261 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 47/50  
 Train metrics: Loss: 0.5296 - Acc: 19.7834 - AUC: 0.5764  
 Valid metrics: Loss: 0.5232 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 48/50  
 Train metrics: Loss: 0.5361 - Acc: 20.2708 - AUC: 0.5404  
 Valid metrics: Loss: 0.5263 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 49/50  
 Train metrics: Loss: 0.5188 - Acc: 20.1805 - AUC: 0.6419  
 Valid metrics: Loss: 0.5232 - Acc: 16.5217 - AUC: 0.1667  
 Epoch 50/50  
 Train metrics: Loss: 0.5218 - Acc: 20.0433 - AUC: 0.5028  
 Valid metrics: Loss: 0.5254 - Acc: 16.5217 - AUC: 0.1667

```
plot(fit_nn)
```



Based on the graphs, we can see that the accuracy on the test and train data are both low (with test accuracy ending up lower). Also, both are relatively steady at the beginning but decrease as the number of epochs gets higher. For the AUC value, both the train and test values end around the same value, but the train AUC value is more steady towards the end of the epochs while the test AUC jumps around. Finally, for the loss, both the test and train data decrease a significant amount of value, once again with the training loss having a more steady/consistent decrease while the test loss fluctuates more.